

PCA and SVD

```

X = df.drop('fake', axis=1)
y = df['fake']

# Perform PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

# Perform SVD
svd = TruncatedSVD(n_components=2)
X_svd = svd.fit_transform(X)

# Print the explained variance ratios
print("Explained variance ratio (PCA):", pca.explained_variance_ratio_)
print("Explained variance ratio (SVD):", svd.explained_variance_ratio_)

Explained variance ratio (PCA): [9.99998806e-01 1.01968536e-06]
Explained variance ratio (SVD): [9.99998803e-01 1.02129951e-06]

X_train, X_test, y_train, y_test = train_test_split(df.drop('fake', axis=1), df['fake'], test_size=0.2, random_state=42)

print('X_train : ')
print(X_train.head())
print('')
print('X_test : ')
print(X_test.head())
print('')
print('y_train : ')
print(y_train.head())
print('')
print('y_test : ')
print(y_test.head())

```

	profile	pic	nums/length	username	fullname	words	nums/length	fullname	\
437		0		0.4		1		0.0	
63		1		0.0		1		0.0	
208		1		0.0		0		0.0	
60		1		0.0		1		0.0	
15		1		0.0		3		0.0	

	name==username	description	length	external	URL	private	#posts	\
437		0		0		1	0	
63		0		68		1	0	100
208		0		18		0	1	133
60		0		3		0	0	12
15		0		46		0	0	254

	#followers	#follows
437	24	88
63	802	151
208	1008	517
60	108	97
15	50374	900


```

X_test :

```

	profile	pic	nums/length	username	fullname	words	nums/length	fullname	\
234		1		0.17		2		0.00	
118		1		0.00		0		0.00	
346		1		0.31		1		0.31	
498		0		0.28		1		0.24	
402		1		0.00		2		0.00	

	name==username	description	length	external	URL	private	#posts	\
234		0		134		1	0	31
118		0		123		1	0	107
346		1		0		0	1	0
498		0		0		0	0	0
402		0		0		0	0	14

	#followers	#follows
234	265	321
118	971	2047
346	25	86
498	86	0
402	143	500


```

y_train :

```

437	1
63	0
208	0
60	0

```
name: fake, dtype: int64

y_test :
234    0
118    0
346    1
498    1
402    1
Name: fake, dtype: int64

#scaling the values into the range 0-1
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df= pd.DataFrame(scaler.fit_transform(df))
df
```

	0	1	2	3	4	5	6	7	8	9	10	11
0	1.0	0.293478	0.000000	0.00	0.0	0.353333	0.0	0.0	0.004331	0.000065	0.127333	0.0
1	1.0	0.000000	0.166667	0.00	0.0	0.293333	0.0	0.0	0.038706	0.000179	0.071067	0.0
2	1.0	0.108696	0.166667	0.00	0.0	0.000000	0.0	1.0	0.001759	0.000010	0.013067	0.0
3	1.0	0.000000	0.083333	0.00	0.0	0.546667	0.0	0.0	0.091893	0.000027	0.086800	0.0
4	1.0	0.000000	0.166667	0.00	0.0	0.000000	0.0	1.0	0.000812	0.000010	0.016800	0.0
...
571	1.0	0.597826	0.083333	0.44	0.0	0.000000	0.0	0.0	0.004466	0.000011	0.079467	1.0
572	1.0	0.413043	0.083333	0.33	0.0	0.140000	0.0	0.0	0.005955	0.000004	0.010000	1.0
573	1.0	0.619565	0.166667	0.00	0.0	0.000000	0.0	0.0	0.000541	0.000006	0.045200	1.0
574	1.0	0.619565	0.083333	0.00	0.0	0.073333	0.0	0.0	0.000000	0.000004	0.009733	1.0
575	1.0	0.293478	0.083333	0.00	0.0	0.000000	0.0	0.0	0.000271	0.000010	0.064933	1.0

576 rows x 12 columns

Decision Tree

```
#building the Decision Tree classifier
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
model = dt.fit(X_train, y_train)

from sklearn.tree import DecisionTreeRegressor

reg = DecisionTreeRegressor(max_depth=2)

reg.fit(X, y)
DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=2,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, #min_impurity_split=None,

min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, #presort='deprecated',
random_state=None, splitter='best')
```

▼

DecisionTreeRegressor

DecisionTreeRegressor(criterion='mse', max_depth=2)

```
from sklearn.tree import plot_tree

plot_tree(reg)
```

```

[Text(0.5, 0.75, 'x[7] <= 0.5\nsquared_error = 0.213\nsamples = 120\nvalue = 0.308'),
Text(0.25, 0.25, 'squared_error = 0.0\nsamples = 83\nvalue = 0.0'),
Text(0.75, 0.25, 'squared_error = 0.0\nsamples = 37\nvalue = 1.0')]

y_hat = dt.predict(X_test)

|squared_error = 0.213|

from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_hat)

0.8706896551724138

|squared_error = 0.0| |squared_error = 0.0|

Logistic Regression

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)

/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
  LogisticRegression
LogisticRegression()

# predicting the test set results
y_pred = model.predict(X_test)

from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
print('R-squared value:', r2)

R-squared value: 0.6178496555855046

```

Linear Regression

```

# create linear regression model
from sklearn import linear_model, metrics
model = linear_model.LinearRegression()

model.fit(X_train, y_train)

  LinearRegression
LinearRegression()

print('Coefficients: \n', model.coef_)

Coefficients:
[-4.48562952e-01  7.97309015e-01 -3.62284700e-02 -9.56923547e-02
 2.36095048e-01 -1.69006009e-03 -1.31382930e-01  9.56927289e-03
-8.01725424e-05 -1.06252148e-08 -2.00946860e-05]

# predicting the test set results
y_pred = model.predict(X_test)

# Results of Linear Regression.
from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_test, y_pred)
print("Mean Square Error : ", mse)

Mean Square Error :  0.11467750277173532

from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
print('R-squared value:', r2)

R-squared value: 0.5378555024568823

```

Logistic Regression